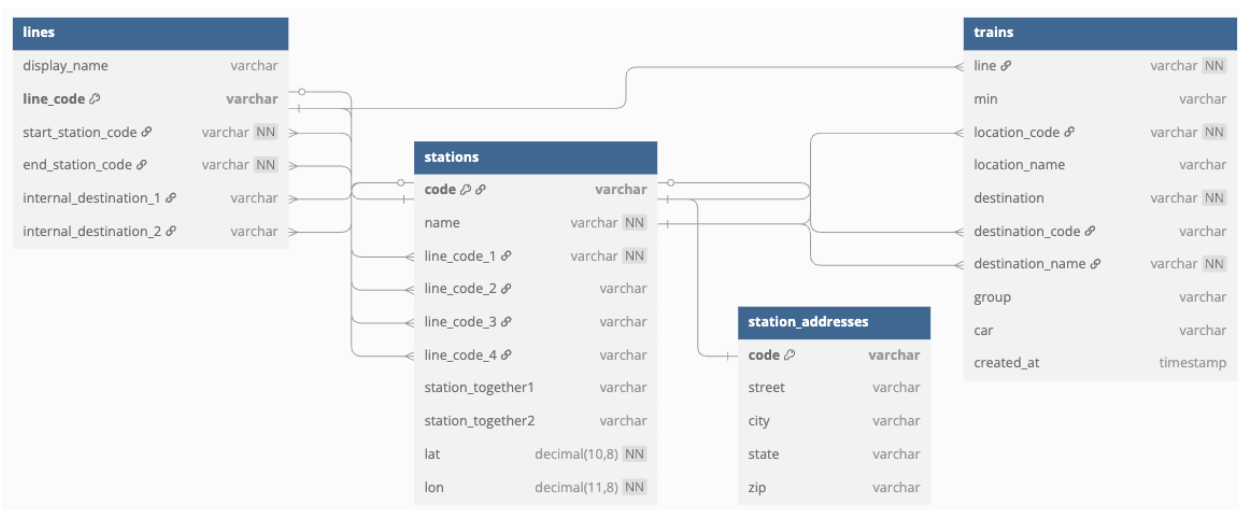


Metro Train Prediction App - Database Architecture Guide

Author: David Morrison

Project Repo: <https://github.com/DavMorr/wmata-app>



DBML: <https://dbdiagram.io/d/WMAT-API-data-mapping-v2-684071dc76955641c2a87aa8>

- Metro Train Prediction App Database Architecture Guide 1
 - Overview 3
 - Database Design Philosophy 3
 - Supported Use Cases..... 4
 - Database Schema..... 5
 - Entity Relationship Diagram..... 5
 - Table Summary..... 6
 - Entity Relationships 6
 - Lines to Stations (Many-to-Many) 6
 - Lines to Station Paths (One-to-Many) 6
 - Stations to Station Addresses (One-to-One)..... 7
 - Connected Stations (Self-Referencing) 7
 - Table Specifications 7
 - lines..... 7

stations.....	8
station_addresses	9
station_paths	9
Indexing Strategy.....	10
Primary Indexes	10
Performance Indexes	11
stations Table.....	11
station_paths Table	11
station_addresses Table	11
Index Usage Analysis.....	11
Query Patterns.....	12
Common Application Queries.....	12
1. Get All Lines for Selection	12
2. Get Ordered Stations for Line	12
3. Get Stations Serving Multiple Lines (Transfer Stations)	12
4. Get Station with Address Information.....	12
5. Find Nearby Stations (Geospatial)	13
Administrative Queries	13
Data Sync Operations	13
Cache Warming Queries	13
Performance Optimization.....	14
Query Performance Targets	14
Optimization Strategies	14
1. Denormalization for Performance.....	14
2. Composite Index Usage	14
3. String Primary Key Optimization.....	15
4. Geospatial Query Optimization	15
Database Configuration Recommendations	15
MySQL Settings	15

Index Monitoring	15
Data Integrity	16
Referential Integrity Constraints	16
Foreign Key Relationships	16
Unique Constraints	16
Data Validation Rules	17
Application-Level Validations	17
Database Constraints	17
Data Consistency Checks	17
Station-Line Relationship Validation	17
Sequence Consistency Validation	18
Migration Strategy	18
Migration Execution Order	18
Migration Files	19
1. Create Lines Table	19
2. Create Stations Table	19
Rollback Strategy	20
Data Seeding Strategy	1

Overview

Database Design Philosophy

The Metro Train Prediction App database is designed around the real-world complexity of the Washington Metro system, emphasizing:

- **Natural Primary Keys** - Using actual Metro codes (line codes, station codes) as primary keys
- **Transfer Station Support** - Multi-line stations handled through flexible line code fields
- **Geographic Ordering** - Station sequences maintained for proper route display
- **High-Precision Geospatial Data** - Coordinates stored with sub-meter accuracy
- **Performance-First Indexing** - Indexes designed around actual query patterns

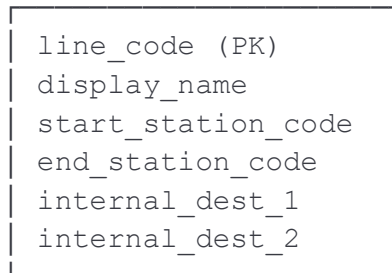
Supported Use Cases

- Progressive form navigation (Lines → Stations → Predictions)
- Transfer station handling (stations serving multiple lines)
- Geographic route planning and distance calculations
- Real-time prediction display with station metadata
- Administrative data synchronization from WMATA API

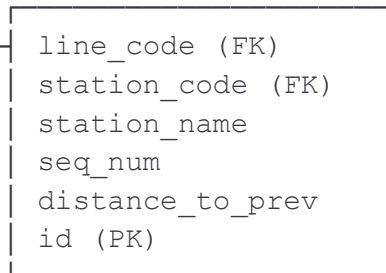
Database Schema

Entity Relationship Diagram

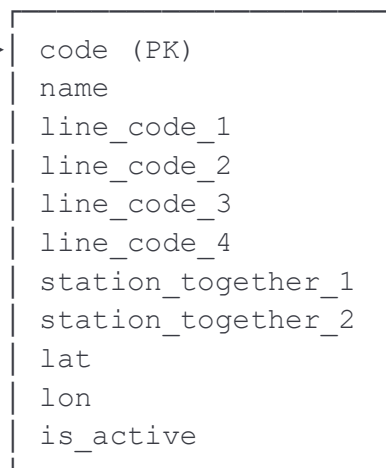
LINES



STATION_PATHS



STATIONS



1:1

STATION_ADDRESSES

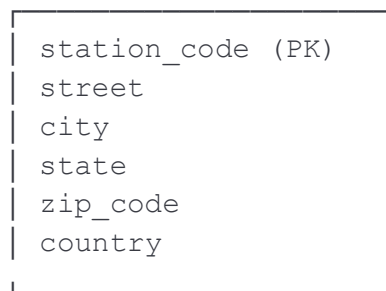


Table Summary

Table	Purpose	Primary Key	Records
lines	Metro line definitions	line_code	~6
stations	Station master data	code	~95
station_addresses	Station location details	station_code	~95
station_paths	Geographic station ordering	id	~95

Entity Relationships

Lines to Stations (Many-to-Many)

Relationship: A line serves multiple stations; a station can serve multiple lines

Implementation: Through `line_code_1`, `line_code_2`, `line_code_3`, `line_code_4` fields in stations table

Examples:

sql

```
-- Red Line serves stations A15 through B11
```

```
SELECT * FROM stations WHERE line_code_1 = 'RD' OR line_code_2 = 'RD';
```

```
-- Metro Center (A01) serves Red, Blue, Orange, Silver Lines
```

```
SELECT * FROM stations WHERE code = 'A01';
```

```
-- Result: line_code_1='RD', line_code_2='BL', line_code_3='OR',  
line_code_4='SV'
```

Lines to Station Paths (One-to-Many)

Relationship: Each line has an ordered sequence of stations

Implementation: `station_paths` table with `line_code` foreign key and `seq_num` ordering

Purpose: Provides geographic ordering for station lists in the frontend

Stations to Station Addresses (One-to-One)

Relationship: Each station has exactly one address

Implementation: `station_addresses` table with `station_code` primary key matching `stations.code`

Cascade Behavior: Deleting a station removes its address

Connected Stations (Self-Referencing)

Relationship: Stations can reference other station platforms

Implementation: `station_together_1` and `station_together_2` fields

Example: Metro Center has platforms A01 and C01 (connected but separate codes)

Table Specifications

lines

Purpose: Metro line definitions and termination points

sql

```
CREATE TABLE lines (  
    line_code VARCHAR(2) PRIMARY KEY,           -- RD, BL, GR, OR, SV, YL  
    display_name VARCHAR(50) NOT NULL,           -- Red, Blue, Green, etc.  
    start_station_code VARCHAR(3) NOT NULL,      -- A15 (Shady Grove)  
    end_station_code VARCHAR(3) NOT NULL,        -- B11 (Glenmont)  
    internal_destination_1 VARCHAR(3),           -- Branch/split destinations  
    internal_destination_2 VARCHAR(3),  
    created_at TIMESTAMP,  
    updated_at TIMESTAMP,  
  
    INDEX idx_display_name (display_name)  
);
```

Key Features:

- Natural primary key using WMATA line codes
- References to start/end stations for route definition
- Internal destinations for handling route branches
- Display name index for alphabetical sorting

Sample Data:

```
sql
INSERT INTO lines VALUES
('RD', 'Red', 'A15', 'B11', NULL, NULL),
('BL', 'Blue', 'J03', 'G05', NULL, NULL),
('GR', 'Green', 'E10', 'F11', 'C15', 'D13');
```

stations

Purpose: Master station data with multi-line support

```
sql
CREATE TABLE stations (
    code VARCHAR(3) PRIMARY KEY,           -- A01, B02, C03, etc.
    name VARCHAR(100) NOT NULL,           -- Metro Center, Union Station
    line_code_1 VARCHAR(2),               -- Primary Line
    line_code_2 VARCHAR(2),               -- Transfer Line 1
    line_code_3 VARCHAR(2),               -- Transfer Line 2
    line_code_4 VARCHAR(2),               -- Transfer Line 3
    station_together_1 VARCHAR(3),         -- Connected platform 1
    station_together_2 VARCHAR(3),         -- Connected platform 2
    lat DECIMAL(10,8) NOT NULL,           -- 38.89834567
    lon DECIMAL(11,8) NOT NULL,           -- -77.02834567
    is_active BOOLEAN DEFAULT true,
    created_at TIMESTAMP,
    updated_at TIMESTAMP,

    INDEX idx_name (name),
    INDEX idx_coordinates (lat, lon),
    INDEX idx_is_active (is_active)
);
```

Key Features:

- 3-character station codes as natural primary keys
- Up to 4 line codes for complex transfer stations
- High-precision coordinates (11mm accuracy)
- Connected station references for platform relationships
- Active status for filtering operational stations

Sample Data:

sql

```
INSERT INTO stations VALUES
('A01', 'Metro Center', 'RD', 'BL', 'OR', 'SV', 'C01', NULL, 38.89834567, -
77.02834567, true),
('A02', 'Farragut North', 'RD', NULL, NULL, NULL, NULL, NULL, 38.90344123, -
77.03927456, true),
('L01', 'Gallery Pl-Chinatown', 'RD', 'YL', 'GR', NULL, 'B01', 'F01',
38.89766789, -77.02112345, true);
```

station_addresses

Purpose: Physical address information for stations

sql

```
CREATE TABLE station_addresses (
    station_code VARCHAR(3) PRIMARY KEY,           -- Matches stations.code
    street VARCHAR(255) NOT NULL,                   -- 1001 G St NW
    city VARCHAR(100) NOT NULL,                     -- Washington
    state VARCHAR(2) NOT NULL,                      -- DC
    zip_code VARCHAR(10) NOT NULL,                   -- 20001-1234
    country VARCHAR(2) DEFAULT 'US',
    created_at TIMESTAMP,
    updated_at TIMESTAMP,

    FOREIGN KEY (station_code) REFERENCES stations(code) ON DELETE CASCADE,
    INDEX idx_city_state (city, state),
    INDEX idx_zip_code (zip_code)
);
```

Key Features:

- One-to-one relationship with stations
- Cascade delete maintains referential integrity
- Indexes for location-based searches
- Supports international addresses (country field)

station_paths

Purpose: Geographic ordering of stations along metro lines

sql

```
CREATE TABLE station_paths (
    id BIGINT PRIMARY KEY AUTO_INCREMENT,
```

```

line_code VARCHAR(2) NOT NULL,           -- RD, BL, etc.
station_code VARCHAR(3) NOT NULL,        -- A01, B02, etc.
station_name VARCHAR(100) NOT NULL,      -- Duplicated for performance
seq_num INTEGER NOT NULL,               -- 1, 2, 3... (geographic
order)
distance_to_prev INTEGER DEFAULT 0,      -- Meters to previous station
created_at TIMESTAMP,
updated_at TIMESTAMP,

FOREIGN KEY (line_code) REFERENCES lines(line_code),
FOREIGN KEY (station_code) REFERENCES stations(code),
INDEX idx_line_sequence (line_code, seq_num),
INDEX idx_station_code (station_code),
UNIQUE KEY uk_line_station (line_code, station_code)
);

```

Key Features:

- Surrogate primary key for flexibility
- Composite foreign keys to both lines and stations
- Sequential numbering for geographic ordering
- Distance tracking for travel time calculations
- Unique constraint prevents duplicate entries

Sample Data:

```

sql
INSERT INTO station_paths VALUES
(1, 'RD', 'A15', 'Shady Grove', 1, 0),
(2, 'RD', 'A14', 'Rockville', 2, 4823),
(3, 'RD', 'A13', 'Twinbrook', 3, 2134),
(4, 'RD', 'A12', 'White Flint', 4, 1876);

```

Indexing Strategy

Primary Indexes

All tables use optimized primary keys:

- `lines.line_code` - String primary key (2 chars)
- `stations.code` - String primary key (3 chars)

- `station_addresses.station_code` - String primary key (3 chars)
- `station_paths.id` - Auto-increment integer

Performance Indexes

stations Table

```
sql
INDEX idx_name (name)           -- Station name searches/sorting
INDEX idx_coordinates (lat, lon) -- Geospatial queries
INDEX idx_is_active (is_active) -- Filter active stations
```

station_paths Table

```
sql
INDEX idx_line_sequence (line_code, seq_num) -- Ordered station retrieval
INDEX idx_station_code (station_code)        -- Individual station lookup
UNIQUE KEY uk_line_station (line_code, station_code) -- Prevent duplicates
```

station_addresses Table

```
sql
INDEX idx_city_state (city, state) -- Location-based searches
INDEX idx_zip_code (zip_code)      -- Postal code lookups
```

Index Usage Analysis

```
sql
-- This query uses idx_line_sequence for optimal performance
SELECT * FROM station_paths
WHERE line_code = 'RD'
ORDER BY seq_num;

-- This query uses idx_coordinates for geospatial searches
SELECT * FROM stations
WHERE lat BETWEEN 38.89 AND 38.91
AND lon BETWEEN -77.04 AND -77.02;

-- Multi-line station query (no specific index, but fast due to small table)
SELECT * FROM stations
WHERE line_code_1 = 'RD' OR line_code_2 = 'RD'
OR line_code_3 = 'RD' OR line_code_4 = 'RD';
```

Query Patterns

Common Application Queries

1. Get All Lines for Selection

sql

```
SELECT line_code as value, display_name as label
FROM lines
ORDER BY display_name;
```

Performance: ~1ms (6 rows, primary key scan)

2. Get Ordered Stations for Line

sql

```
SELECT sp.station_code as value, sp.station_name as label,
       sp.seq_num, sp.distance_to_prev
FROM station_paths sp
WHERE sp.line_code = ?
ORDER BY sp.seq_num;
```

Performance: ~5ms (uses idx_line_sequence)

3. Get Stations Serving Multiple Lines (Transfer Stations)

sql

```
SELECT s.code, s.name, s.line_code_1, s.line_code_2, s.line_code_3,
       s.line_code_4
FROM stations s
WHERE s.line_code_1 = ? OR s.line_code_2 = ?
      OR s.line_code_3 = ? OR s.line_code_4 = ?;
```

Performance: ~10ms (table scan, but small table)

4. Get Station with Address Information

sql

```
SELECT s.code, s.name, s.lat, s.lon,
       sa.street, sa.city, sa.state, sa.zip_code
FROM stations s
LEFT JOIN station_addresses sa ON s.code = sa.station_code
WHERE s.code = ?;
```

Performance: ~2ms (primary key lookup + join)

5. Find Nearby Stations (Geospatial)

sql

```
SELECT code, name, lat, lon,
       (6371000 * acos(cos(radians(?)) * cos(radians(lat)) *
        cos(radians(lon) - radians(?)) + sin(radians(?)) *
        sin(radians(lat)))) AS distance
FROM stations
WHERE lat BETWEEN ? - 0.01 AND ? + 0.01
      AND lon BETWEEN ? - 0.01 AND ? + 0.01
      AND is_active = true
ORDER BY distance
LIMIT 10;
```

Performance: ~15ms (uses idx_coordinates for initial filtering)

Administrative Queries

Data Sync Operations

sql

-- Insert/Update Lines

```
INSERT INTO lines (line_code, display_name, start_station_code,
end_station_code)
VALUES (?, ?, ?, ?)
ON DUPLICATE KEY UPDATE
display_name = VALUES(display_name),
start_station_code = VALUES(start_station_code),
end_station_code = VALUES(end_station_code);
```

-- Clear and Rebuild Paths

```
DELETE FROM station_paths WHERE line_code = ?;
INSERT INTO station_paths (line_code, station_code, station_name, seq_num,
distance_to_prev)
VALUES (?, ?, ?, ?, ?);
```

Cache Warming Queries

sql

-- Preload all station-line relationships

```
SELECT s.code, s.name, s.line_code_1, s.line_code_2, s.line_code_3,
s.line_code_4
FROM stations s
WHERE s.is_active = true;
```

```
-- Preload all line paths
SELECT line_code, station_code, station_name, seq_num
FROM station_paths
ORDER BY line_code, seq_num;
```

Performance Optimization

Query Performance Targets

- **Line selection:** < 5ms
- **Station list for line:** < 10ms
- **Station details:** < 5ms
- **Geospatial searches:** < 20ms
- **Multi-line station queries:** < 15ms

Optimization Strategies

1. Denormalization for Performance

The `station_paths` table duplicates station names for performance:

```
sql
-- Instead of always joining to stations table
SELECT sp.station_code, s.name
FROM station_paths sp
JOIN stations s ON sp.station_code = s.code
WHERE sp.line_code = 'RD';
```

```
-- We can query directly (faster)
SELECT station_code, station_name
FROM station_paths
WHERE line_code = 'RD';
```

2. Composite Index Usage

The `(line_code, seq_num)` index optimizes the most common query pattern:

```
sql
EXPLAIN SELECT * FROM station_paths WHERE line_code = 'RD' ORDER BY seq_num;
-- Uses: idx_line_sequence (covering index)
```

```
-- Rows examined: ~27 (only Red Line stations)
-- Extra: Using index
```

3. String Primary Key Optimization

Using actual Metro codes as primary keys provides benefits:

- **Meaningful joins:** `WHERE line_code = 'RD'` is more readable than `WHERE line_id = 1`
- **Cache efficiency:** Natural keys are more cache-friendly
- **Reduced joins:** No need to join for display values

4. Geospatial Query Optimization

For location-based queries, use bounding box filtering before distance calculations:

```
sql
-- Efficient: Filter with index first, then calculate distance
SELECT *, (complex_distance_calculation) as distance
FROM stations
WHERE lat BETWEEN ? AND ?      -- Uses index for initial filtering
      AND lon BETWEEN ? AND ?  -- Further reduces candidate set
      AND is_active = true
ORDER BY distance;
```

Database Configuration Recommendations

MySQL Settings

```
sql
-- Optimize for read-heavy workload
innodb_buffer_pool_size = 1G          -- Cache frequently accessed data
query_cache_size = 128M               -- Cache repeated queries
query_cache_type = 1                  -- Enable query cache

-- Index optimization
innodb_stats_on_metadata = 0          -- Reduce metadata overhead
optimizer_search_depth = 4            -- Optimize join planning
```

Index Monitoring

```
sql
-- Check index usage
SELECT
```

```

        s.table_name,
        s.index_name,
        s.cardinality,
        round(((s.cardinality / t.table_rows) * 100), 2) as selectivity
FROM information_schema.statistics s
INNER JOIN information_schema.tables t
    ON s.table_schema = t.table_schema
    AND s.table_name = t.table_name
WHERE s.table_schema = 'metro_transit'
ORDER BY s.table_name, selectivity;

```

Data Integrity

Referential Integrity Constraints

Foreign Key Relationships

sql

```

-- Station addresses must reference valid stations
ALTER TABLE station_addresses
ADD CONSTRAINT fk_station_addresses_station_code
FOREIGN KEY (station_code) REFERENCES stations(code) ON DELETE CASCADE;

```

```

-- Station paths must reference valid lines and stations
ALTER TABLE station_paths
ADD CONSTRAINT fk_station_paths_line_code
FOREIGN KEY (line_code) REFERENCES lines(line_code);

```

```

ALTER TABLE station_paths
ADD CONSTRAINT fk_station_paths_station_code
FOREIGN KEY (station_code) REFERENCES stations(code);

```

Unique Constraints

sql

```

-- Prevent duplicate station entries per line
ALTER TABLE station_paths
ADD CONSTRAINT uk_line_station UNIQUE (line_code, station_code);

```


Data Validation Rules

Application-Level Validations

php

// Line codes must be 2 uppercase letters

'line_code' => 'required|regex:/^[A-Z]{2}\$/'

// Station codes must be 3 alphanumeric characters

'code' => 'required|regex:/^[A-Z0-9]{3}\$/'

// Coordinates must be within reasonable bounds

'lat' => 'required|numeric|between:38.0,39.5' *// DC area bounds*

'lon' => 'required|numeric|between:-78.0,-76.0'

// Sequence numbers must be positive

'seq_num' => 'required|integer|min:1'

Database Constraints

sql

-- Coordinate bounds checking

ALTER TABLE stations

ADD CONSTRAINT chk_lat_bounds CHECK (lat BETWEEN 38.0 AND 39.5);

ALTER TABLE stations

ADD CONSTRAINT chk_lon_bounds CHECK (lon BETWEEN -78.0 AND -76.0);

-- Sequence numbers must be positive

ALTER TABLE station_paths

ADD CONSTRAINT chk_seq_num_positive CHECK (seq_num > 0);

-- Distance must be non-negative

ALTER TABLE station_paths

ADD CONSTRAINT chk_distance_non_negative CHECK (distance_to_prev >= 0);

Data Consistency Checks

Station-Line Relationship Validation

sql

```

-- Find stations without any line assignments
SELECT code, name
FROM stations
WHERE line_code_1 IS NULL
      AND line_code_2 IS NULL
      AND line_code_3 IS NULL
      AND line_code_4 IS NULL;

-- Find paths referencing non-existent stations
SELECT sp.line_code, sp.station_code
FROM station_paths sp
LEFT JOIN stations s ON sp.station_code = s.code
WHERE s.code IS NULL;

```

Sequence Consistency Validation

sql

```

-- Check for missing sequence numbers
SELECT line_code,
       COUNT(*) as station_count,
       MAX(seq_num) as max_seq,
       MIN(seq_num) as min_seq
FROM station_paths
GROUP BY line_code
HAVING station_count != (max_seq - min_seq + 1);

-- Find duplicate sequence numbers within lines
SELECT line_code, seq_num, COUNT(*)
FROM station_paths
GROUP BY line_code, seq_num
HAVING COUNT(*) > 1;

```

Migration Strategy

Migration Execution Order

1. **Create base tables** (lines, stations)
2. **Create dependent tables** (station_addresses, station_paths)
3. **Add foreign key constraints**
4. **Create indexes**

5. Load initial data

Migration Files

1. Create Lines Table

php

// 2025_06_04_184151_create_lines_table.php

```
Schema::create('lines', function (Blueprint $table) {
    $table->string('line_code', 2)->primary();
    $table->string('display_name', 50);
    $table->string('start_station_code', 3);
    $table->string('end_station_code', 3);
    $table->string('internal_destination_1', 3)->nullable();
    $table->string('internal_destination_2', 3)->nullable();
    $table->timestamps();

    $table->index('display_name');
});
```

2. Create Stations Table

php

// 2025_06_04_184201_create_stations_table.php

```
Schema::create('stations', function (Blueprint $table) {
    $table->string('code', 3)->primary();
    $table->string('name', 100);
    $table->string('line_code_1', 2)->nullable();
    $table->string('line_code_2', 2)->nullable();
    $table->string('line_code_3', 2)->nullable();
    $table->string('line_code_4', 2)->nullable();
    $table->string('station_together_1', 3)->nullable();
    $table->string('station_together_2', 3)->nullable();
    $table->decimal('lat', 10, 8);
    $table->decimal('lon', 11, 8);
    $table->boolean('is_active')->default(true);
    $table->timestamps();

    $table->index('name');
    $table->index(['lat', 'lon']);
    $table->index('is_active');
});
```

Rollback Strategy

php

```
// ALL migrations include proper down() methods
public function down(): void
{
    Schema::dropIfExists('station_paths'); // Drop dependent tables first
    Schema::dropIfExists('station_addresses');
    Schema::dropIfExists('stations');
    Schema::dropIfExists('lines');          // Drop parent tables last
}
```

Data Seeding Strategy

php

```
// DatabaseSeeder.php

public function run(): void {
    // Metro Train Prediction App uses a custom sync command
    // that pulls data directly from the WMATA API
    try {
        $this->command->info('Syncing Metro data from WMATA API...');
        $exitCode = Artisan::call('metro:sync');
        if ($exitCode === 0) {
            $this->command->info('Metro data sync completed successfully');
        }
        else {
            $this->command->error('Metro data sync failed with exit code: ' . $exitCode);
            throw new \Exception('Metro sync command failed');
        }
    }
    catch (\Exception $e) {
        $this->command->error('Failed to sync Metro data: ' . $e->getMessage());
        $this->command->warn('Database seeding will continue, but Metro tables may be empty');
        $this->command->warn('Run "sail artisan metro:sync" manually after resolving API issues');
        // Don't throw - allow other seeders to run
        // throw $e; // Uncomment to fail entire seeding process
    }
}
```

Alternative Seeding Methods

bash

```
# Direct command execution (recommended)
sail artisan metro:sync

# With validation check first
sail artisan metro:sync --validate
```

```
# Via DatabaseSeeder during fresh installation  
sail artisan db:seed
```

The Metro Train Prediction App uses a specialized sync command rather than traditional Laravel seeders because:

- Data is sourced from the live WMATA API
- Station relationships and paths are calculated dynamically
- Ensures data consistency with the current Metro system
- Handles complex multi-line station assignments automatically

This database architecture provides a robust foundation for the Metro Train Prediction App, optimized for the specific query patterns of the frontend while maintaining data integrity and supporting future expansion.